

Linear regression

Ngô Minh Nhựt



2021

Linear regression with one variable

Part 1


Machine learning

- Machine learning is about to find out structure of observed data or relationship inside them




Source: Andrew Ng

Supervised learning



- Learn: provided with input and corresponding output
- Inference: infer output for new input

Linear regression



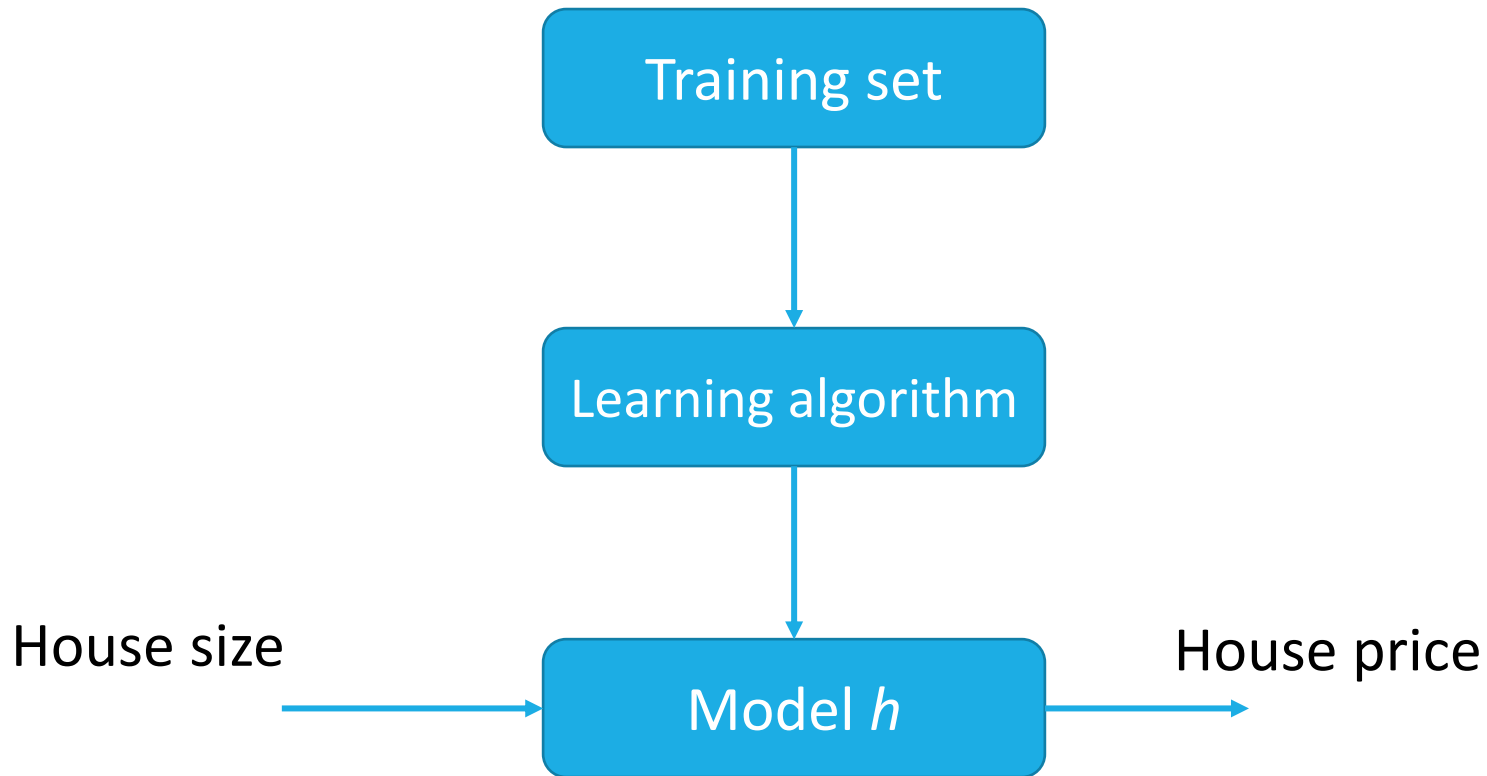
- Learn: provided with input and corresponding output
- Inference: infer output for new input
- **Output: continuous real value**

Training set

House price
with respect to
size

x: Size (m2)	y: Price (millions VND)
20	600
50	876
80	1800
100	2000
...	...

- m : number of samples
- x : input
- y : output/label
- (x, y) : training sample
- $(x^{(i)}, y^{(i)})$: i^{th} sample



- ❑ h : mapping from size to price
- ❑ Linear regression: linear mapping
- ❑ In this part: linear regression with one variable

Hypothesis

- Training set: $(x^{(i)}, y^{(i)}), i = 1, 2, \dots, m$
- Hypothesis : $h(x) = \theta_0 + \theta_1 x$
 - x : input
 - $h(x)$: output
- Goal: identify θ_0 and θ_1 so that model $h(x)$ fits with training set most
 - Given x , identify $h(x)$, so that $h(x)$ closes to y most
 - x : input
 - $h(x)$: estimated output
 - y : realistic output

Hypothesis

House price
with regard to
size

x: Size (m2)	y: Price (millions VND)
20	600
50	876
80	1800
100	2000
...	...

Hypothesis : $h(x) = \theta_0 + \theta_1 x$


x: input

$h(x)$: output

θ_0, θ_1 : parameters

Learning: find out θ_0, θ_1 from training set

Hypothesis



Cost function

- Model error with one sample

- $\frac{1}{2}(h(x) - y)^2 = \frac{1}{2}(\theta_0 + \theta_1 x - y)^2$

- Cost function

- $J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h(x^{(i)}) - y^{(i)})^2$

- $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$

- Goal:

- Identify θ_0, θ_1 so that $J(\theta_0, \theta_1)$ reaches minimal

Cost function

- Hypothesis : $h_{\theta}(x) = \theta_0 + \theta_1 x$
- Parameters: θ_0, θ_1
- Cost function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



- Goal: find out θ_0, θ_1 so that $J(\theta_0, \theta_1)$ reaches minimal

- ❖ $\text{Minimize}_{\theta_0, \theta_1} J(\theta_0, \theta_1)$



- To demonstrate: $\theta_0 = 0, h_{\theta}(x) = \theta_1 x$

- ❖ $J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$



Cost function





Cost function





Cost function




Cost function



Cost function



Cost function



Source: Andrew Ng

Gradient descent algorithm

- Cost function

- $J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h(x^{(i)}) - y^{(i)})^2$

- Goal

- Identify θ_0, θ_1 so that $J(\theta_0, \theta_1)$ reaches minimal

Gradient descent algorithm

- Cost function

- $J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h(x^{(i)}) - y^{(i)})^2$

- Goal

- Identify θ_0, θ_1 so that $J(\theta_0, \theta_1)$ reaches minimal

- To find θ_0, θ_1

- Starting from certain point θ_0, θ_1 (e.g., $\theta_0 = 0, \theta_1 = 0$)
 - Change values of θ_0, θ_1 until $J(\theta_0, \theta_1)$ reaches minimal

- How to change values of θ_0, θ_1 ?

Gradient descent algorithm

□ Partial derivative

- $\frac{dJ}{d\theta_0} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$
- $\frac{dJ}{d\theta_1} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$

Gradient descent algorithm

□ Partial derivative

- $\frac{dJ}{d\theta_0} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$
- $\frac{dJ}{d\theta_1} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$

□ Gradient vector

□ $\left(\frac{dJ}{d\theta_0}, \frac{dJ}{d\theta_1} \right)$

Gradient descent algorithm

□ Partial derivative

- $\frac{dJ}{d\theta_0} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$
- $\frac{dJ}{d\theta_1} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$

□ Gradient vector

□ $\left(\frac{dJ}{d\theta_0}, \frac{dJ}{d\theta_1} \right)$

□ Going backward gradient vector makes function decrease

Gradient descent algorithm

Loop until convergence

{

$$\theta_0 = \theta_0 - \alpha \frac{dJ}{d\theta_0}$$

$$\theta_1 = \theta_1 - \alpha \frac{dJ}{d\theta_1}$$

}


- α is learning rate

- Steps

- Step 1: calculate gradient vector
- Step 2: update elements of vector θ
- Step 3: recalculate cost

Learning rate

- ❑ If learning rate α is too small, algorithm converges slowly
- ❑ If learning rate α is too large, algorithm may not converge



Source: Andrew Ng

Multivariate linear regression

Part 2

Training set

x: Size (m2)	y: Price (millions VND)
20	600
50	876
80	1800
100	2000
...	...

$$\text{Hypothesis : } h_{\theta}(x) = \theta_0 + \theta_1 x$$

Training set

x_1 : Size (m ²)	x_2 : Age (year)	x_3 : Number of floor (m)	y : Price (millions VND)
20	5	1	600
20	8	3	876
80	10	3	1800
70	7	5	2000
...

- $(x^{(i)}, y^{(i)})$: i^{th} sample
- $x^{(i)}$: input of i^{th} sample
- $x_j^{(i)}$: j^{th} feature i^{th} sample
- n : number of features

Hypothesis

- Single variable

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

← single value

- Multiple variables

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

← vector

Hypothesis

□ Multiple variables

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

$$x \in \mathbb{R}^n, \quad \theta \in \mathbb{R}^{n+1}$$

Given $x_0 = 1$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix},$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix}$$

$$\rightarrow h_{\theta}(x) = \theta^T x, \quad \theta, x \in \mathbb{R}^{n+1}$$

Cost function

- Hypothesis

- $h_{\theta}(x) = \theta^T x$

- Parameters

- $\theta = [\theta_0, \theta_1, \theta_2, \dots, \theta_n]^T$

- Cost function

- $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)})^2$

Gradient descent algorithm

- Gradient vector

- $\frac{dJ}{d\theta_j} = \frac{1}{m} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)}) x_j^{(i)}$

- $j = 0, 1, 2, \dots, n$

- Repeat until convergence

{

$$\theta_j = \theta_j - \alpha \frac{dJ}{d\theta_j}$$


}

Feature normalization

x_1 : Size (m ²)	x_2 : Age (year)	x_3 : Number of floor (m)	y: Price (millions VND)
20	5	1	600
20	8	3	876
80	10	3	1800
70	7	5	2000
...

Feature normalization

Scale values of features into the same range



$$x_1 = 20, 30, \dots, 100$$

$$x_2 = 1, 2, \dots, 10$$

$$x_1 = \frac{x_1}{100}$$

$$x_2 = \frac{x_2}{10}$$

Source: Andrew Ng

Feature normalization

- Mean-based normalization


$$x_j = \frac{x_j - \mu_j}{s_j}$$

- μ_j : mean

- s_j : standard deviation

- Range after normalization: $-1 \leq x_j \leq 1$



Learning rate



- Check if J decreases after each step of updating
- J converges when J decreases an amount smaller than 0.001 (ϵ) after a step

Source: Andrew Ng

Learning rate



- With large learning rate, J may not converge
- With small learning rate, J converges slowly
- Try with: ..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, ...

Source: Andrew Ng