

Overfitting, regularization, and model validation




Ngô Minh Nhật

2021

Content

- ❑ Overfitting
- ❑ Regularization
- ❑ Model validation

Overfitting




When there are many features:

- ▶ Model may overfit with training data (training cost ~ 0)
- ▶ Could not generalize for new samples (testing cost $\gg 0$)

Source: Andrew Ng

Overfitting

□ Logistic regression



Source: Andrew Ng

Overfitting

□ Solutions:


■ Reduce number of features

- Select features through data analysis
- Select features through model validation

■ Regularization


- Reduce effect of parameters to output
- Good in case features contributes slightly to output

Regularization



Size of house

$$\theta_0 + \theta_1 x + \theta_2 x^2$$



Size of house

$$\theta_0 + \theta_1 x + \theta_2 x^2 + \cancel{\theta_3 x^3} + \cancel{\theta_4 x^4}$$

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + 1000\theta_3^2 + 1000\theta_4^2$$

Source: Andrew Ng

Regularization

- When values of parameters become small
 - Produce simpler models
 - Reduce overfitting

Regularization

□ Cost function


$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

- Lambda: weight decay
- When lambda is bigger, model becomes simpler

Regularization

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

If lambda too big?



Source: Andrew Ng

Regularization

□ Vector gradient

$$\frac{dJ}{d\theta_0} = \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_0^{(i)}$$

$$\frac{dJ}{d\theta_j} = \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)} + \frac{\lambda}{m} \theta_j$$

Gradient descent algorithm works as usual

Regularization

□ Logistic regression – cost function

$$J(\theta) = - \left[\frac{1}{m} \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \\ + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Regularization

- Logistic regression – vector gradient

$$\frac{dJ}{d\theta_0} = \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_0^{(i)}$$

$$\frac{dJ}{d\theta_j} = \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)} + \frac{\lambda}{m} \theta_j$$


Model validation

- ❑ Learning process makes model fit with training data
- ❑ Can learnt model generalize for new samples?
 - Validate the model with unseen data
- ❑ Solution:
 - Learn, validate, and select the best model
 - Test if the selected model works well with new data



Dataset

- ❑ Split dataset into 3 subsets: train, validation, and test
- ❑ With large dataset, the ratio is: 60% : 20% : 20%



Cost function

Train error

$$J_{\text{train}}(\theta) = \frac{1}{2m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

Cross-validation error

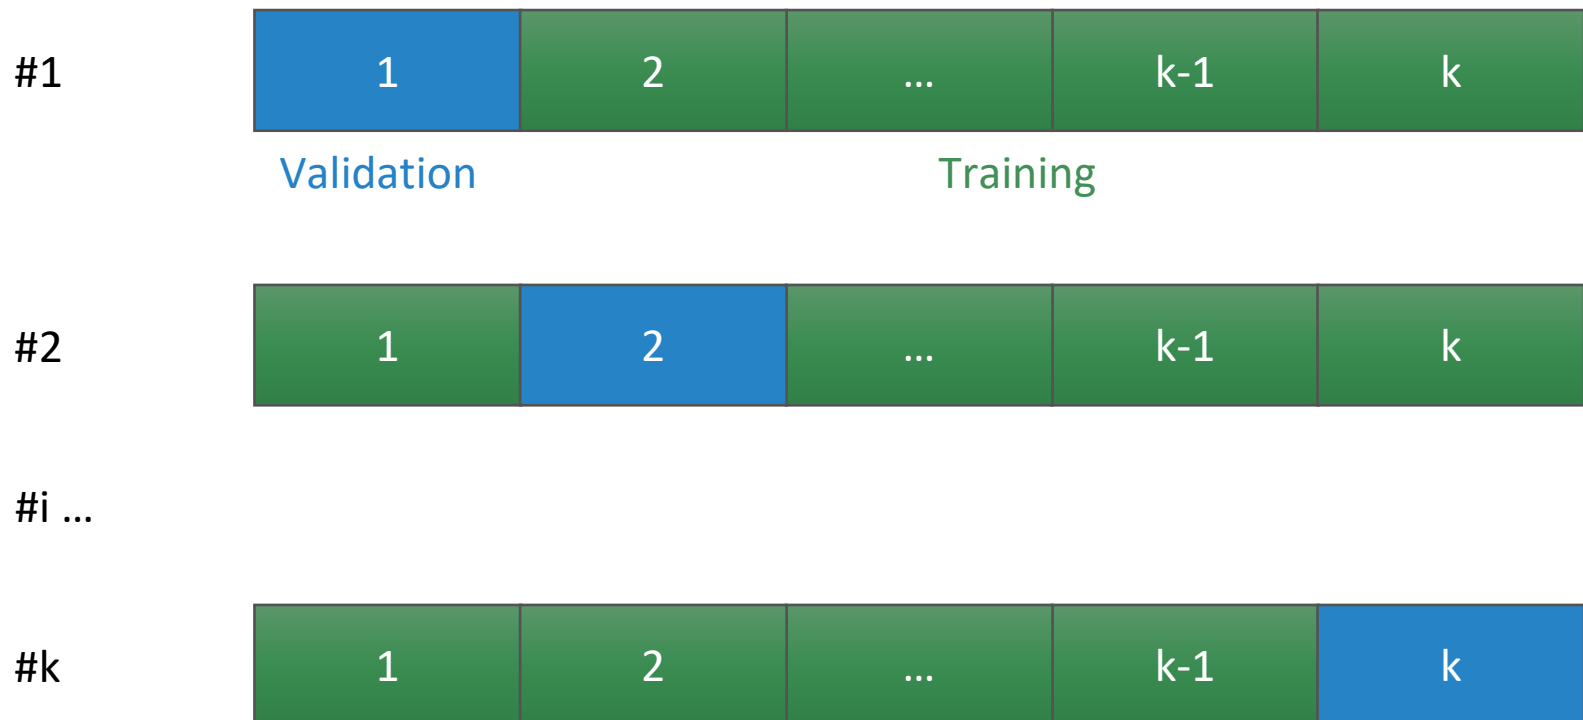
$$J_{\text{cv}}(\theta) = \frac{1}{2m_{\text{cv}}} \sum_{i=1}^{\text{cv}} \left(h_{\theta}(x_{\text{cv}}^{(i)}) - y_{\text{cv}}^{(i)} \right)^2$$

Test error

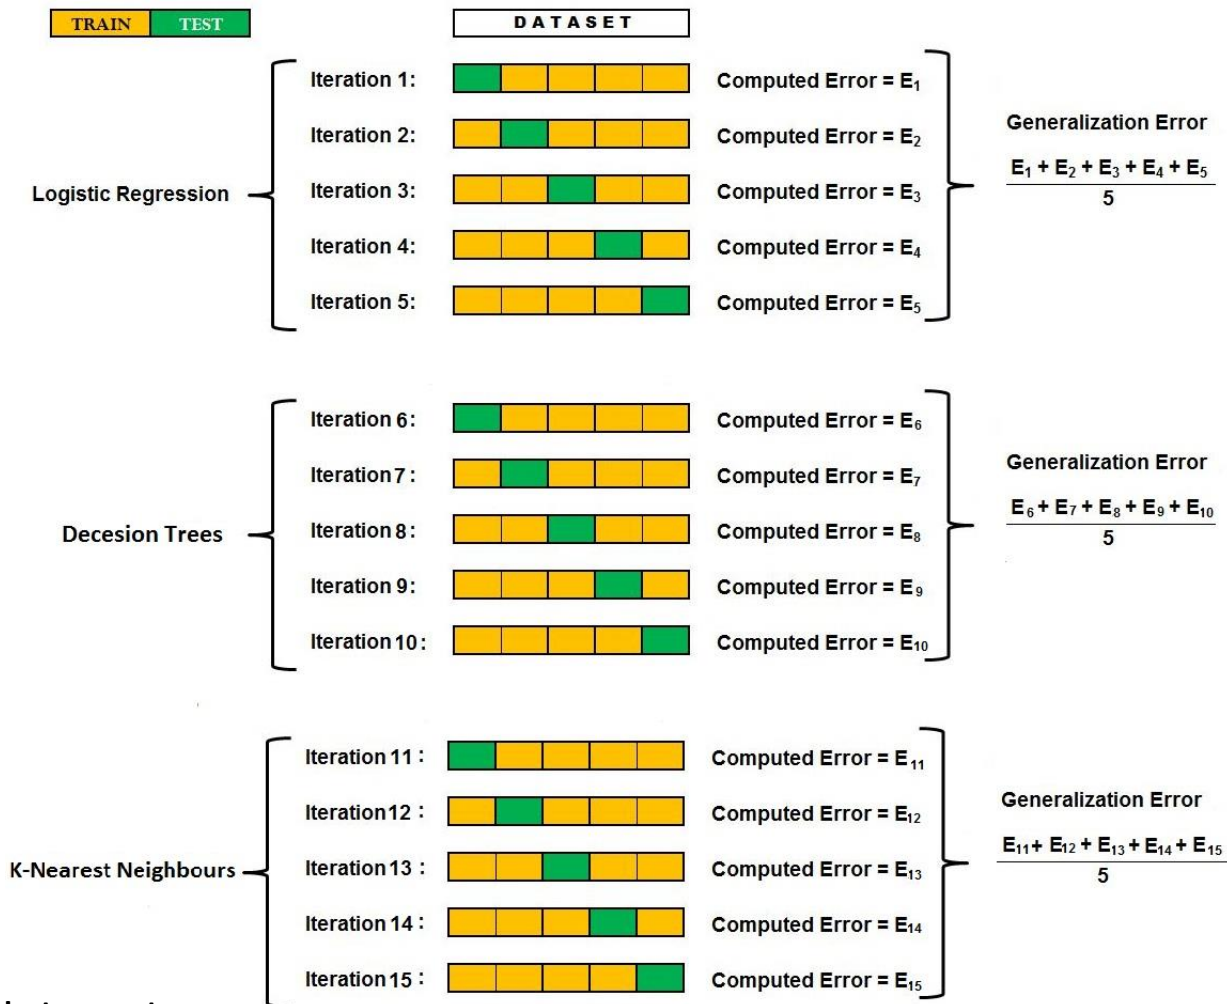
$$J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_{i=1}^{\text{test}} \left(h_{\theta}(x_{\text{test}}^{(i)}) - y_{\text{test}}^{(i)} \right)^2$$

k-fold cross validation

- ❑ Randomly split dataset into k parts
- ❑ Learn and validate k time. k is usually 10



k-fold cross validation




Source: Internet

Model finetuning


- ❑ Model has various parameters
- ❑ Learning algorithm has various hyper parameters
- ❑ What should we do when model does not work well?
 - Collect more data
 - Reduce number of features
 - Try with new features
 - Switch to other models or hypotheses
 - Reduce weight decay
 - Increase weight decay

Model finetuning




Bias and variance

- ❑ Bias: model error on training set
- ❑ Variance: model error on validation set
- ❑ Adjust bias and variance until they reach minimal




$$\theta_0 + \theta_1 x$$

High bias
(underfit)



$$\theta_0 + \theta_1 x + \theta_2 x^2$$


“Just right”



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

High variance
(overfit)

Bias and variance



Model evaluation

- ❑ Measurements
 - Precision, recall
 - Accuracy
 - F1-score
- ❑ Depending on problems, we need some suitable measures


Model evaluation

$$\square \text{ Precision} = \frac{\text{True positive}}{\text{Predicted positive}}$$

$$\square \text{ Recall} = \frac{\text{True positive}}{\text{Positive}}$$

$$\square \text{ F1 - score} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\square \text{ Accuracy} = \frac{\text{True positive} + \text{True negative}}{\text{Positive} + \text{Negative}}$$



How many retrieved items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are retrieved?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Source: Wikipedia

Confusion matrix

		Predicted				Total
		Cat	Dog	Tiger	Wolf	
Actual	Cat	6	0	3	1	10
	Dog	2	4	0	4	10
	Tiger	3	3	3	0	9
	Wolf	1	4	1	2	8
Total		12	11	7	7	

```
>>> sklearn.metrics.classification_report
```

	precision	recall	f1-score	support
Cat	0.500	0.600	0.545	10
Dog	0.364	0.400	0.381	10
Tiger	0.429	0.333	0.375	9
Wolf	0.286	0.250	0.267	8
accuracy			0.405	37
macro avg	0.394	0.396	0.392	37
weighted avg	0.399	0.405	0.399	37

```
>>> y_true = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]
>>> y_pred = [0, 0, 0, 0, 0, 0, 2, 2, 2, 3, ...]
>>> target_names = ['Cat', 'Dog', 'Tiger', 'Wolf']
```